

.ipynb

October 25, 2019

```
[1]: import numpy as np
import scipy.optimize
import random
import csv

[2]: ### Tasks - Regression

### 1. What is the distribution of ratings in the dataset? That is,
### how many 1-star, 2-star, 3-star (etc.) reviews are there?
### You may write out the values or include a simple plot (1 mark).

# Read the tsv file and extract the rating star, verified status
# and length of review as three different lists.
data = []
with open('amazon_reviews_us_Gift_Card_v1_00.tsv') as tsvfile:
    reader = csv.reader(tsvfile, delimiter='\t')
    for row in reader:
        data.append(row)

#print(len(data))
#data = np.array(data)
headline = data[0]
#print(headline)

rating = []
purStatus = []
lenOfReview = []
for i in range(len(data)):
    rating.append(data[i][7])
    purStatus.append(data[i][11])
    lenOfReview.append(len(data[i][13]))
#print(len(rating))
#print(len(purStatus))
#print(len(lenOfReview))

[3]: # check the distribution of stars by count function
numOf1Star = rating.count('1')
```

```

numOf2Star = rating.count('2')
numOf3Star = rating.count('3')
numOf4Star = rating.count('4')
numOf5Star = rating.count('5')

```

```

[4]: # print out the total count for star number 1,2,3 relatively
print('Total number of 1 star: ' , numOf1Star)
print('Total number of 2 star: ' , numOf2Star)
print('Total number of 3 star: ' , numOf3Star)

```

```

Total number of 1 star: 4766
Total number of 2 star: 1560
Total number of 3 star: 3147

```

```

[5]: # trim the verified status data
verRv = purStatus[1:]
#print(verRv[0])
for i in range(len(verRv)):
    if verRv[i] == 'Y':
        verRv[i] = 1
    else: verRv[i] = 0
#print(len(verRv))
#print(verRv[0:5])

```

```

[6]: # trim the length of review data
lenOfReview = lenOfReview[1:]

```

```

[7]: # set up our X matrix for the following regression problem
colOf1 = [1]*len(verRv)
X = np.column_stack((colOf1,verRv,lenOfReview))
X = y = np.asmatrix(X,dtype='float')
#X.shape

```

```

[8]: # set up our y vector for the following regression problem
y = rating[1:]
y = np.asmatrix(y,dtype='float')
#y.shape

```

```

[9]: # conduct the regression progress by the function
theta,residuals,rank,s = np.linalg.lstsq(X,y.T)

```

```

//anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:2: FutureWarning:
`rcond` parameter will change to the default of machine precision times ``max(M,
N)`` where M and N are the input matrix dimensions.
To use the future default and silence this warning we advise to pass
`rcond=None`, to keep using the old, explicitly pass `rcond=-1`.

```

```
[10]: ### 3. Train a simple predictor to predict the star rating using two features
### Report the values of 0, 1, and 2. Briefly describe your interpretation of
→these values
### Explain these in terms of the features and labels

# check the intercept and coefficients. Theta 0 is the coefficient, theta 1
# is the coefficient for verified status, theta 2 is the coefficient of review
# length.
print('Theta0 = ',theta[0])
print('Theta1 = ',theta[1])
print('Theta2 = ',theta[2])
```

```
Theta0 = [[4.84503504]]
Theta1 = [[0.04985806]]
Theta2 = [[-0.00124546]]
```

```
[11]: #Theta 0 represents the mean value of our prediction when the purchase is not
→verified and the review length is 0
#Theta 1 represents the mean difference of the same length review when the
→purchase is verified or not
#Theta 2 represents the mean increment comes with one increment of review
→length when the verified status is fixed

# The coefficients theta1 is positive. Therefore, if the review length is
→fixed, the verified purchase will have a
# higher average rating number with 0.0498. Samely, the review length
→coefficient is with a negative theta -0.001245.
# That means with a purchase in a fixed purchase status, we have an average
→decrement of 0.001245 on the rating numbers.
# That also means a longer review tends to have a negative effect on our final
→rating stars. So, a negative review is
# likely to be longer than a positive one.
```

```
[12]: np.linalg.inv(X.T * X) * X.T * y.T
```

```
[12]: matrix([[ 4.84503504e+00],
              [ 4.98580589e-02],
              [-1.24545526e-03]])
```

```
[13]: ### 4. Train another predictor that only uses one feature:
X1 = X[:,0:2]
X1.shape
```

```
[13]: (148310, 2)
```

```
[14]: theta1,residuals,rans,s = np.linalg.lstsq(X1,y.T)
```

```
//anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:1: FutureWarning:
`rcond` parameter will change to the default of machine precision times ``max(M,
```

N)` where M and N are the input matrix dimensions.
 To use the future default and silence this warning we advise to pass
 `rcond=None`, to keep using the old, explicitly pass `rcond=-1`.
 """Entry point for launching an IPython kernel.

```
[15]: ### Report the values of 0 and 1
print('Theta0 = ',theta1[0])
print('Theta1 = ',theta1[1])
```

```
Theta0 = [[4.578143]]
Theta1 = [[0.16793392]]
```

```
[16]: ### Provide an explanation as to why these coefficients might vary so
      →significantly (1 mark).1

#The coefficients are different because the review length is not under our
      →consideration now.
#The linear regression try to solve them with a minimized MRE, without the
      →review length, the minimization progress
#is quite different now. We are only try to optimize the rating stars towards
      →the verified status. Therefore, the
#coefficients would vary through this progress.
```

```
[17]: ### 5. Split the data into two fractions the first 90% for training, and the
      →remaining 10%
      ### testing (based on the order they appear in the file). Train the same model
      →as in Question
      ### 4 on the training set only. What is the models MSE on the training and on
      →the test set (1 mark)?

# Split the 90% and 10% of our data
X1train = X1[0: round(0.9*(X1.shape[0])),]
print(X1train.shape)
X1test = X1[round(0.9*(X1.shape[0])):,]
print(X1test.shape)
print(y.shape)
ytrain = y[0,0: round(0.9*(y.shape[1]))]
ytest = y[0,round(0.9*(y.shape[1])):]
print(ytrain.shape)
print(ytest.shape)
```

```
(133479, 2)
(14831, 2)
(1, 148310)
(1, 133479)
(1, 14831)
```

```
[18]: thetaT,residualsT,rankT,sT = np.linalg.lstsq(X1train,ytrain.T)
```

```
//anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:1: FutureWarning:  
`rcond` parameter will change to the default of machine precision times ``max(M,  
N)`` where M and N are the input matrix dimensions.  
To use the future default and silence this warning we advise to pass  
`rcond=None`, to keep using the old, explicitly pass `rcond=-1`.  
    """Entry point for launching an IPython kernel.
```

```
[19]: print(residualsT)
```

```
[[87493.37815713]]
```

```
[20]: # Calculate the test MSE  
print(ytest.shape)  
ypredict = X1test@thetaT  
print(ypredict.shape)  
difference = np.subtract(ypredict,ytest.T)  
print(difference.shape)
```

```
(1, 14831)
```

```
(14831, 1)
```

```
(14831, 1)
```

```
[21]: SumtestMSE = np.linalg.norm(difference)  
#print(SumtestMSE)  
testMSE = (SumtestMSE**2)/ypredict.shape[0]  
print('test MSE',testMSE)
```

```
test MSE 0.9723851990303902
```

```
[22]: ### report the MSE for training and test sets  
MSETrain = (residualsT/ytrain.shape[1])[0,0]  
print('MSE for the training residual',MSETrain)  
print('MSE for the testing residual',testMSE)
```

```
MSE for the training residual 0.6554842196685237
```

```
MSE for the testing residual 0.9723851990303902
```

```
[23]: ### 7. (CSE258 only) Repeat the above experiment, varying the size  
### of the training and test fractions between 5% and 95% for training  
### (using the complement for testing). Show how the training and test  
### error vary as a function of the training set size (again using a  
### simple plot or table). Does the size of the training set make a  
### significant difference in testing performance? Comment on why it
```

```
### might or might not make a significant difference in this instance (2 marks).
```

```
# the fuction could change the size of training and test sets
```

```
def separateTrainNTest(X, y, percentage):  
    XTrain = X[0: round(percentage*(X.shape[0])),]  
    XTest = X[round(percentage*(X.shape[0])):,]  
    yTrain = y[0,0: round(percentage*(y.shape[1]))]  
    yTest = y[0,round(percentage*(y.shape[1])):]  
    return [XTrain,XTest,yTrain,yTest]
```

```
[24]: # conduct a new MSE calculator of regression for the training set
```

```
def MSEregressionOfTrain(XTrain,yTrain):  
    thetaTrain,residualsTrain,rankTrain,sTrain = np.linalg.lstsq(XTrain,yTrain.  
→T)  
    MSETrain = (residualsTrain/yTrain.shape[1])[0,0]  
    return [thetaTrain,MSETrain]
```

```
[25]: # conduct a new MSE calculator of regression for the test set
```

```
def MSEregressionOfTest(thetaTrain,XTest,yTest):  
    ypredict = XTest@thetaTrain  
    difference = np.subtract(ypredict,yTest.T)  
    SumtestMSE = np.linalg.norm(difference)  
    MSETest = (SumtestMSE**2)/ypredict.shape[0]  
    return MSETest
```

```
[26]: MSErecord = []
```

```
for i in range(91):  
    XTrain = separateTrainNTest(X1,y,(i+5)*0.01)[0]  
    XTest = separateTrainNTest(X1,y,(i+5)*0.01)[1]  
    yTrain = separateTrainNTest(X1,y,(i+5)*0.01)[2]  
    yTest = separateTrainNTest(X1,y,(i+5)*0.01)[3]  
    MSETrain = MSEregressionOfTrain(XTrain,yTrain)[1]  
    thetaTrain = MSEregressionOfTrain(XTrain,yTrain)[0]  
    MSETest = MSEregressionOfTest(thetaTrain,XTest,yTest)  
    MSErecord.append([MSETrain,MSETest])  
MSErecord = np.asarray(MSErecord)  
print(MSErecord.shape)
```

//anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:3: FutureWarning:
`rcond` parameter will change to the default of machine precision times ``max(M,
N)`` where M and N are the input matrix dimensions.

To use the future default and silence this warning we advise to pass
`rcond=None`, to keep using the old, explicitly pass `rcond=-1`.

This is separate from the ipykernel package so we can avoid doing imports
until

(91, 2)

```
[27]: # plot the MSE values into a plot
import matplotlib.pyplot as plt

xAxis = np.array(range(5,96))
# plot the train MSE values into a plot
plt.plot(xAxis,MSErecord[:,0],'g--')
# plot the test MSE values into a plot
plt.plot(xAxis,MSErecord[:,1], 'o--')
plt.title('Comparison of Train(green)/Test(blue) MSE based on different size')
plt.xlabel('Size of Percentage of Training Set')
plt.ylabel('MSE Value')

# We see that with the size of training set increase, the MSE goes down a
  ↳ little bit, since more sets of
# data means the regression becomes more precise in the first place. However,
  ↳ when the data is large enough
# the regression line would become relatively stable. So, the training MSE goes
  ↳ up a little since more data
# comes with more possible error.
# We see that as the size of training sets increase, the MSE for test increase
  ↳ as well. Since the increase
# of training set size means the MSE optimization more fit the training set.
  ↳ So, the test MSE goes up.
```

[27]: Text(0, 0.5, 'MSE Value')

```
[28]: ### Tasks Classification (week 2):
### 8. First, lets train a predictor that estimates whether a review
### is verified using the rating and the length:
###  $p(\text{review is verified}) = 0 + 1 \cdot \mathbb{E}[\text{star rating}] + 2 \cdot \mathbb{E}[\text{review length}]$ 
### Train a logistic regressor to make the above prediction
### (you may use a logistic regression library with default parameters,
### e.g. linear model.LogisticRegression() from sklearn).
### Report the classification accuracy of this predictor.
### Report also the proportion of labels that are positive
### (i.e., the proportion of reviews that are verified) and
### the proportion of predictions that are positive (1 mark).

# prepare the data for logistic classification
yLogistic = X[:,1].T
#print(yLogistic)
#print(X.shape)
XLogistic = []
XLogistic = np.asarray(XLogistic)
XLogistic = np.column_stack((X[:,0],y.T, X[:,2]))
#print(XLogistic)
# prepare the training data for logistic classification
XLTrain = XLogistic[0: round(0.9*(XLogistic.shape[0])),]
```

```

XLTest = XLogistic[round(0.9*(XLogistic.shape[0])):,:]
print(XLTrain.shape)
print(XLTest.shape)
yLtrain = yLogistic[0,0: round(0.9*(yLogistic.shape[1]))]
yLtest = yLogistic[0,round(0.9*(yLogistic.shape[1])):]
print(yLtrain.shape)
print(yLtest.shape)

```

```

(133479, 3)
(14831, 3)
(1, 133479)
(1, 14831)

```

```
[29]: from sklearn.linear_model import LogisticRegression
```

```
[30]: clf = LogisticRegression(random_state=0, solver='lbfgs',multi_class='ovr').
      →fit(XLTrain, yLtrain.T)
```

```

//anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:724:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
    y = column_or_1d(y, warn=True)

```

```
[31]: clf.predict(XLTest)
print('The proportion of positive labels:',sum(clf.predict(XLTest))/14831)
proportionOfVerified = sum(yLtest.T)/14831
print('The proportion of verified review purchase:',proportionOfVerified[0,0])
clf.score(XLTest, yLtest.T)
print('The accuracy for prediction based on star rating and review length is:
      →',clf.score(XLTest, yLtest.T))

```

```

The proportion of positive labels: 0.9989886049490931
The proportion of verified review purchase: 0.5595711684984155
The accuracy for prediction based on star rating and review length is:
0.5597734475085968

```

```
[32]: from datetime import datetime
date = []
for i in range(len(data)):
    date.append(data[i][-1])
date = date[1:]
#print(date)

fmt = '%Y-%m-%d'
d = datetime.strptime('2014-10-14',fmt) #start date
for i in range(len(date)):

```



```

        diff = datetime.strptime(date[i],fmt) - d
        date[i] = diff.days

# date = [datetime.strptime(x,'%Y-%m-%d') for x in date]
min(date)

```

[32]: -3652

```

[33]: # prepare the data for logistic classification over my own chosen variables
# I choose the review date
yLogistic1 = X[:,1].T
XLogistic1 = []
XLogistic1 = np.asarray(XLogistic1)
XLogistic1 = np.column_stack((X[:,0],date))
#print(XLogistic1)

# prepare the training data for logistic classification

XLTrain1 = XLogistic1[0: round(0.9*(XLogistic1.shape[0])),]
XLTest1 = XLogistic1[round(0.9*(XLogistic1.shape[0])):,]
#print(XLTrain1.shape)
#print(XLTest1.shape)

yLtrain1 = yLogistic1[0,0: round(0.9*(yLogistic1.shape[1]))]
yLtest1 = yLogistic1[0,round(0.9*(yLogistic1.shape[1])):]
#print(yLtrain1.shape)
#print(yLtest1.shape)

```

```

[34]: clf1 = LogisticRegression(random_state=0, solver='lbfgs',multi_class='ovr').
      →fit(XLTrain1, yLtrain1.T)

```

```

//anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:724:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
    y = column_or_1d(y, warn=True)

```

```

[35]: clf1.predict(XLTest1)
sum(clf1.predict(XLTest1))
print('The accuracy for prediction based on review date is:',clf1.
      →score(XLTest1, yLtest1.T))
print('The proportion of positive label prediction based on review date is:
      →',sum(clf1.predict(XLTest1))/14831)
proportionOfVerified1 = sum(yLtest1.T)/14831
print('The proportion of verified review purchase:',proportionOfVerified1[0,0])

```

The accuracy for prediction based on review date is: 0.5638190277122244
The proportion of positive label prediction based on review date is:

0.9907625918683838

The proportion of verified review purchase: 0.5595711684984155